

12: Support Vector Machines (SVMs)

[Previous](#) [Next](#) [Index](#)

Support Vector Machine (SVM) - Optimization objective


- So far, we've seen a range of different algorithms
 - With supervised learning algorithms - performance is pretty similar
 - What matters more often is;
 - The amount of training data
 - Skill of applying algorithms
- One final supervised learning algorithm that is widely used - **support vector machine (SVM)**
 - Compared to both logistic regression and neural networks, a SVM sometimes gives a cleaner way of learning non-linear functions
 - Later in the course we'll do a survey of different supervised learning algorithms

An alternative view of logistic regression

- Start with logistic regression, see how we can modify it to get the SVM
 - As before, the logistic regression hypothesis is as follows


$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

- And the sigmoid activation function looks like this



- In order to explain the math, we use z as defined above
- What do we want logistic regression to do?
 - We have an example where $y = 1$
 - Then we hope $h_{\theta}(x)$ is close to 1

- With $h_{\theta}(x)$ close to 1, $(\theta^T x)$ must be **much larger** than 0




- Similarly, when $y = 0$
 - Then we hope $h_{\theta}(x)$ is close to 0
 - With $h_{\theta}(x)$ close to 0, $(\theta^T x)$ must be **much less** than 0
- This is our classic view of logistic regression
 - Let's consider another way of thinking about the problem
- Alternative view of logistic regression
 - If you look at cost function, each example contributes a term like the one below to the overall cost function

$$-(y \log h_{\theta}(x) + (1 - y) \log(1 - h_{\theta}(x)))$$

- For the overall cost function, we sum over all the training examples using the above function, and have a $1/m$ term
- If you then plug in the hypothesis definition ($h_{\theta}(x)$), you get an expanded cost function equation;


$$= -y \log \frac{1}{1 + e^{-\theta^T x}} - (1 - y) \log(1 - \frac{1}{1 + e^{-\theta^T x}})$$

- So each training example contributes that term to the cost function for logistic regression
- If $y = 1$ then only the first term in the objective matters
 - If we plot the functions vs. z we get the following graph



- This plot shows the cost contribution of an example when $y = 1$ given z
 - So if z is big, the cost is low - this is good!
 - But if z is 0 or negative the cost contribution is high
 - This is why, when logistic regression sees a positive example, it tries to set $\theta^T x$ to be a very large term
- If $y = 0$ then only the second term matters


- We can again plot it and get a similar graph



- Same deal, if z is small then the cost is low
 - But if z is large then the cost is massive


SVM cost functions from logistic regression cost functions

- To build a SVM we must redefine our cost functions
 - When $y = 1$
 - Take the $y = 1$ function and create a new cost function
 - Instead of a curved line create two straight lines (magenta) which acts as an approximation to the logistic regression $y = 1$ function



- Take point (1) on the z axis
 - Flat from 1 onwards
 - Grows when we reach 1 or a lower number
 - This means we have two straight lines
 - Flat when cost is 0
 - Straight growing line after 1
 - So this is the new $y=1$ cost function
 - Gives the SVM a computational advantage and an easier optimization problem
 - We call this function **cost₁(z)**
- Similarly
 - When $y = 0$

- Do the equivalent with the $y=0$ function plot



- We call this function **cost₀(z)**

- So here we define the two cost function terms for our SVM graphically
 - How do we implement this?

The complete SVM cost function

- As a comparison/reminder we have logistic regression below

$$\min_{\theta} \frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \left(-\log h_{\theta}(x^{(i)}) \right) + (1 - y^{(i)}) \left(-\log(1 - h_{\theta}(x^{(i)})) \right) \right] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

- If this looks unfamiliar its because we previously had the - sign outside the expression
- For the SVM we take our two logistic regression $y=1$ and $y=0$ terms described previously and replace with
 - $\text{cost}_1(\theta^T x)$
 - $\text{cost}_0(\theta^T x)$
- So we get

$$\min_{\theta} \frac{1}{m} \sum_{i=1}^m \left[y^{(i)} \text{cost}_1(\theta^T x^{(i)}) + (1 - y^{(i)}) \text{cost}_0(\theta^T x^{(i)}) \right] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

SVM notation is slightly different

- In convention with SVM notation we rename a few things here
- 1) Get rid of the $1/m$ terms
 - This is just a slightly different convention
 - By removing $1/m$ we should get the same optimal values for
 - $1/m$ is a constant, so should get same optimization
 - e.g. say you have a minimization problem which minimizes to $u = 5$
 - If your cost function * by a constant, you still generates the minimal value
 - That minimal value is different, but that's irrelevant
- 2) For logistic regression we had two terms:
 - Training data set term (i.e. that we sum over m) = **A**
 - Regularization term (i.e. that we sum over n) = **B**
 - So we could describe it as $A + \lambda B$
 - Need some way to deal with the trade-off between regularization and data set terms
 - Set different values for λ to parametrize this trade-off
 - Instead of parameterization this as $A + \lambda B$
 - For SVMs the convention is to use a different parameter called C
 - So do $CA + B$

- If C were equal to $1/\lambda$ then the two functions ($CA + B$ and $A + \lambda B$) would give the same value
- So, our overall equation is

$$\min_{\theta} C \sum_{i=1}^m \left[y^{(i)} \text{cost}_1(\theta^T x^{(i)}) + (1 - y^{(i)}) \text{cost}_0(\theta^T x^{(i)}) \right] + \frac{1}{2} \sum_{j=1}^n \theta_j^2$$

- Unlike logistic, $h_{\theta}(x)$ doesn't give us a probability, but instead we get a direct prediction of 1 or 0
 - So if $\theta^T x$ is equal to or greater than 0 $\rightarrow h_{\theta}(x) = 1$
 - Else $\rightarrow h_{\theta}(x) = 0$

Large margin intuition

- Sometimes people refer to SVM as **large margin classifiers**
 - We'll consider what that means and what an SVM hypothesis looks like
 - The SVM cost function is as above, and we've drawn out the cost terms below



If $y = 1$, we want $\theta^T x \geq 1$ (not just ≥ 0)

If $y = 0$, we want $\theta^T x \leq -1$ (not just < 0)

- Left is cost_1 and right is cost_0
- What does it take to make terms small
 - If $y = 1$
 - $\text{cost}_1(z) = 0$ only when $z \geq 1$
 - If $y = 0$
 - $\text{cost}_0(z) = 0$ only when $z \leq -1$
- Interesting property of SVM
 - If you have a positive example, you only really *need* z to be greater or equal to 0
 - If this is the case then you predict 1
 - SVM wants a bit more than that - doesn't want to *just* get it right, but have the value be quite a bit bigger than zero
 - Throws in an extra safety margin factor
- Logistic regression does something similar
- What are the consequences of this?
 - Consider a case where we set C to be huge
 - $C = 100,000$
 - So considering we're minimizing $CA + B$
 - If C is huge we're going to pick an A value so that A is equal to zero
 - What is the optimization problem here - how do we make $A = 0$?
 - Making $A = 0$
 - If $y = 1$
 - Then to make our " A " term 0 need to find a value of θ so $(\theta^T x)$ is greater than or equal to 1
 - Similarly, if $y = 0$
 - Then we want to make " A " = 0 then we need to find a value of θ so $(\theta^T x)$ is equal to or less than -1


- So - if we think of our optimization problem a way to ensure that this first "A" term is equal to 0, we re-factor our optimization problem into just minimizing the "B" (regularization) term, because
 - When $A = 0 \rightarrow A^*C = 0$
- So we're minimizing B, under the constraints shown below

$$\min \frac{1}{2} \sum_{i=1}^n \theta_j^2$$


$$\text{s.t. } \theta^T x^{(i)} \geq 1 \quad \text{if } y^{(i)} = 1$$

$$\theta^T x^{(i)} \leq -1 \quad \text{if } y^{(i)} = 0$$

- Turns out when you solve this problem you get interesting decision boundaries




- The green and magenta lines are functional decision boundaries which could be chosen by logistic regression
 - But they probably don't generalize too well
- The black line, by contrast is the chosen by the SVM because of this safety net imposed by the optimization graph
 - More robust separator
- Mathematically, that black line has a larger minimum distance (margin) from any of the training examples




- By separating with the largest margin you incorporate robustness into your decision making process
- We looked at this at when C is very large
 - SVM is more sophisticated than the large margin might look

- If you were just using large margin then SVM would be very sensitive to outliers



- You would risk making a ridiculous hugely impact your classification boundary
 - A single example might not represent a good reason to change an algorithm
 - If C is very large then we *do* use this quite naive maximize the margin approach



- So we'd change the black to the magenta
 - But if C is reasonably small, or a not too large, then you stick with the black decision boundary
- What about non-linearly separable data?
 - Then SVM still does the right thing if you use a normal size C
 - So the idea of SVM being a large margin classifier is only really relevant when you have no outliers and you can easily linearly separable data
- Means we ignore a few outliers

Large margin classification mathematics (optional)


Vector inner products

- Have two (2D) vectors u and v - what is the inner product $(u^T v)$?


$$u = \begin{bmatrix} u_1 \\ u_2 \end{bmatrix} \quad v = \begin{bmatrix} v_1 \\ v_2 \end{bmatrix}$$

- Plot u on graph

- i.e u_1 vs. u_2




- One property which is good to have is the **norm** of a vector
 - Written as $||u||$
 - This is the euclidean length of vector u
 - So $||u|| = \text{SQRT}(u_1^2 + u_2^2) = \text{real number}$
 - i.e. length of the arrow above
 - Can show via Pythagoras
- For the inner product, take v and orthogonally project down onto u
 - First we can plot v on the same axis in the same way (v_1 vs v_2)
 - Measure the length/magnitude of the projection



- So here, the green line is the projection
 - p = length along u to the intersection
 - p is the magnitude of the projection of vector v onto vector u
- Possible to show that
 - $u^T v = p * ||u||$
 - So this is one way to compute the inner product
 - $u^T v = u_1 v_1 + u_2 v_2$
 - So therefore
 - **$p * ||u|| = u_1 v_1 + u_2 v_2$**
 - This is an important rule in linear algebra
 - We can reverse this too
 - So we could do
 - $v^T u = v_1 u_1 + v_2 u_2$
 - Which would obviously give you the same number

- p can be negative if the angle between them is 90 degrees or more



- So here p is negative
- Use the vector inner product theory to try and understand SVMs a little better

SVM decision boundary

$$\min_{\theta} \frac{1}{2} \sum_{j=1}^n \theta_j^2$$

$$\text{s.t. } \theta^T x^{(i)} \geq 1 \quad \text{if } y^{(i)} = 1$$

$$\theta^T x^{(i)} \leq -1 \quad \text{if } y^{(i)} = 0$$

- For the following explanation - two simplification
 - Set $\theta_0 = 0$ (i.e. ignore intercept terms)
 - Set $n = 2$ - (x_1, x_2)
 - i.e. each example has only 2 features
- Given we only have two parameters we can simplify our function to

$$\frac{1}{2} (\theta_1^2 + \theta_2^2)$$

- And, can be re-written as

$$\frac{1}{2} \left(\sqrt{\theta_1^2 + \theta_2^2} \right)^2$$

- Should give same thing
- We may notice that

$$\frac{1}{2} \left(\sqrt{\theta_1^2 + \theta_2^2} \right)^2$$


$$= ||\theta||$$

- The term in red is the norm of θ
 - If we take θ as a 2×1 vector
 - If we assume $\theta_0 = 0$ its still true
- So, finally, this means our optimization function can be re-defined as


$$= \frac{1}{2} ||\theta||^2$$

- So the SVM is minimizing the squared norm
- Given this, what are the $(\theta^T x)$ parameters doing?
 - Given θ and given example x what is this equal to
 - We can look at this in a comparable manner to how we just looked at u and v


- Say we have a single positive training example (red cross below)




- Although we haven't been thinking about examples as vectors it can be described as such



- Now, say we have our parameter vector θ and we plot that on the same axis



- The next question is what is the inner product of these two vectors



- p , is in fact p^i , because it's the length of p for example i
 - Given our previous discussion we know

$$(\theta^T x^i) = p^i * ||\theta||$$

$$= \theta_1 x_1^i + \theta_2 x_2^i$$

- So these are both equally valid ways of computing $\theta^T x^i$
- What does this mean?
 - The constraints we defined earlier
 - $(\theta^T x) \geq 1$ if $y = 1$
 - $(\theta^T x) \leq -1$ if $y = 0$
 - Can be replaced/substituted with the constraints
 - $p^i \cdot \|\theta\| \geq 1$ if $y = 1$
 - $p^i \cdot \|\theta\| \leq -1$ if $y = 0$

- Writing that into our optimization objective

$$\min_{\theta} \frac{1}{2} \sum_{j=1}^n \theta_j^2 = \frac{1}{2} \|\theta\|^2$$


$$\text{s.t. } p^{(i)} \cdot \|\theta\| \geq 1 \quad \text{if } y^{(i)} = 1$$

$$p^{(i)} \cdot \|\theta\| \leq -1 \quad \text{if } y^{(i)} = 0$$


- So, given we've redefined these functions let us now consider the training example below




- Given this data, what boundary will the SVM choose? Note that we're still assuming $\theta_0 = 0$, which means the boundary has to pass through the origin (0,0)
 - Green line - small margins




- SVM would not choose this line
 - Decision boundary comes very close to examples
 - Lets discuss *why* the SVM would **not** choose this decision boundary
 - Looking at this line
 - We can show that θ is at 90 degrees to the decision boundary




- **θ is always at 90 degrees to the decision boundary** (can show with linear algebra, although we're not going to!)
- So now let's look at what this implies for the optimization objective
 - Look at first example (x^1)




- Project a line from x^1 on to to the θ vector (so it hits at 90 degrees)
 - The distance between the intersection and the origin is (p^1)
 - Similarly, look at second example (x^2)
 - Project a line from x^2 into to the θ vector
 - This is the magenta line, which will be **negative** (p^2)
 - If we overview these two lines below we see a graphical representation of what's going on;




- We find that both these p values are going to be pretty small
 - If we look back at our optimization objective
 - We know we need $p^1 * ||\theta||$ to be bigger than or equal to 1 for positive examples
 - If p is small
 - Means that $||\theta||$ must be pretty large
 - Similarly, for negative examples we need $p^2 * ||\theta||$ to be smaller than or equal to -1
 - We saw in this example p^2 is a small negative number
 - So $||\theta||$ must be a large number
 - Why is this a problem?
 - The optimization objective is trying to find a set of parameters where the norm of theta is small
 - So this doesn't seem like a good direction for the parameter vector (because as p values get smaller $||\theta||$ must get larger to compensate)
 - So we should make p values larger which allows $||\theta||$ to become smaller
- So let's choose a different boundary



- Now if you look at the projection of the examples to θ we find that p^1 becomes large and $||\theta||$ can become small
- So with some values drawn in



- This means that by choosing this second decision boundary we can make $||\theta||$ smaller
 - Which is why the SVM chooses this hypothesis as better
 - This is how we generate the large margin effect




- The magnitude of this margin is a function of the p values
 - So by maximizing these p values we minimize $||\theta||$
- Finally, we did this derivation assuming $\theta_0 = 0$,
 - If this is the case we're entertaining only decision boundaries which pass through $(0,0)$
 - If you allow θ_0 to be other values then this simply means you can have decision boundaries which cross through the x and y values at points other than $(0,0)$
 - Can show with basically same logic that this works, and even when θ_0 is non-zero when you have optimization objective described above (when C is very large) that the SVM is looking for a large margin separator between the classes


Kernels - 1: Adapting SVM to non-linear classifiers

- What are kernels and how do we use them
 - We have a training set

- We want to find a non-linear boundary



- Come up with a complex set of polynomial features to fit the data
 - Have $h_\theta(x)$ which
 - Returns 1 if the combined weighted sum of vectors (weighted by the parameter vector) is less than or equal to 0
 - Else return 0
 - Another way of writing this (new notation) is
 - That a hypothesis computes a decision boundary by taking the sum of the parameter vector multiplied by a **new feature vector f** , which simply contains the various high order x terms
 - e.g.
 - $h_\theta(x) = \theta_0 + \theta_1 f_1 + \theta_2 f_2 + \theta_3 f_3$
 - Where
 - $f_1 = x_1$
 - $f_2 = x_1 x_2$
 - $f_3 = \dots$
 - i.e. not specific values, but each of the terms from your complex polynomial function
 - Is there a better choice of feature f than the high order polynomials?
 - As we saw with computer imaging, high order polynomials become computationally expensive
- New features
 - Define three features in this example (ignore x_0)
 - Have a graph of x_1 vs. x_2 (don't plot the values, just define the space)
 - Pick three points in that space



- These points l^1 , l^2 , and l^3 , were chosen manually and are called **landmarks**
 - Given x , define f_1 as the similarity between (x, l^1)

$$= \exp(-(\|x - l^1\|^2) / 2\sigma^2)$$

$$\exp\left(-\frac{\|x - l^{(1)}\|^2}{2\sigma^2}\right)$$

- $\|x - l^1\|$ is the euclidean distance between the point x and the landmark l^1 squared
 - Discussed more later
- If we remember our statistics, we know that
 - σ is the **standard deviation**
 - σ^2 is commonly called the **variance**
- Remember, that as discussed

$$\|x - l^{(1)}\|^2 = \sum_{j=1}^n (x_j - l_j^{(1)})^2$$


- So, f_2 is defined as
 - $f_2 = \text{similarity}(x, l^1) = \exp(-(\|x - l^1\|^2) / 2\sigma^2)$
- And similarly
 - $f_3 = \text{similarity}(x, l^2) = \exp(-(\|x - l^2\|^2) / 2\sigma^2)$
- This similarity function is called a **kernel**
 - This function is a **Gaussian Kernel**
- So, instead of writing similarity between x and l we might write
 - $f_1 = k(x, l^1)$

Diving deeper into the kernel

- So let's see what these kernels do and why the functions defined make sense
 - Say x is close to a landmark
 - Then the squared distance will be ~ 0
 - So


$$f_1 \approx \exp\left(-\frac{0^2}{2\sigma^2}\right)$$

- Which is basically e^{-0}
 - Which is close to 1
 - Say x is far from a landmark
 - Then the squared distance is big
 - Gives $e^{-\text{large number}}$
 - Which is close to zero
 - Each landmark defines a new feature
- If we plot f_1 vs the kernel function we get a plot like this
 - Notice that when $x = [3, 5]$ then $f_1 = 1$
 - As x moves away from $[3, 5]$ then the feature takes on values close to zero
 - So this measures how close x is to this landmark




What does σ do?

- σ^2 is a parameter of the Gaussian kernel
 - Defines the steepness of the rise around the landmark
- Above example $\sigma^2 = 1$
- Below $\sigma^2 = 0.5$




◦ We see here that as you move away from 3,5 the feature f_1 falls to zero much more rapidly

- The inverse can be seen if $\sigma^2 = 3$




- Given this definition, what kinds of hypotheses can we learn?
 - With training examples x we predict "1" when
 - $\theta_0 + \theta_1 f_1 + \theta_2 f_2 + \theta_3 f_3 \geq 0$
 - For our example, let's say we've already run an algorithm and got the
 - $\theta_0 = -0.5$
 - $\theta_1 = 1$
 - $\theta_2 = 1$
 - $\theta_3 = 0$
 - Given our placement of three examples, what happens if we evaluate an example at the **magenta dot** below?




- Looking at our formula, we know f_1 will be close to 1, but f_2 and f_3 will be close to 0
 - So if we look at the formula we have
 - $\theta_0 + \theta_1 f_1 + \theta_2 f_2 + \theta_3 f_3 \geq 0$
 - $-0.5 + 1 + 0 + 0 = 0.5$
 - 0.5 is greater than 0

- If we had **another point** far away from all three



- This equates to -0.5
 - So we predict 0
- Considering our parameter, for points near l^1 and l^2 you predict 1, but for points near l^3 you predict 0
- Which means we create a non-linear decision boundary that goes a lil' something like this;



- Inside we predict $y = 1$
- Outside we predict $y = 0$
- So this show how we can create a non-linear boundary with landmarks and the kernel function in the support vector machine
 - But
 - How do we get/chose the landmarks
 - What other kernels can we use (other than the Gaussian kernel)

Kernels II

- Filling in missing detail and practical implications regarding kernels
- Spoke about picking landmarks manually, defining the kernel, and building a hypothesis function
 - Where do we get the landmarks from?
 - For complex problems we probably want lots of them

Choosing the landmarks

- Take the training data
- For each example place a landmark at exactly the same location
- So end up with m landmarks
 - One landmark per location per training example
 - Means our features measure how close to a training set example something is
- Given a new example, compute all the f values
 - Gives you a feature vector f (f_0 to f_m)
 - $f_0 = 1$ always
- A more detailed look at generating the f vector
 - If we had a training example - features we compute would be using (x^i, y^i)

- So we just cycle through each landmark, calculating how close to that landmark actually x^i is
 - $f_1^i = k(x^i, l^1)$
 - $f_2^i = k(x^i, l^2)$
 - ...
 - $f_m^i = k(x^i, l^m)$
- Somewhere in the list we compare x to itself... (i.e. when we're at f_1^i)
 - So because we're using the Gaussian Kernel this evaluates to 1
- Take these m features ($f_1, f_2 \dots f_m$) group them into an $[m + 1 \times 1]$ dimensional vector called f
 - f^i is the f feature vector for the i th example
 - And add a 0th term = 1
- Given these kernels, how do we use a support vector machine

SVM hypothesis prediction with kernels

- Predict $y = 1$ if $(\theta^T f) \geq 0$
 - Because $\theta = [m+1 \times 1]$
 - And $f = [m + 1 \times 1]$
- So, this is how you make a prediction assuming you already have θ
 - How do you get θ ?

SVM training with kernels

- Use the SVM learning algorithm

$$\min_{\theta} C \sum_{i=1}^m y^{(i)} \text{cost}_1(\theta^T f^{(i)}) + (1 - y^{(i)}) \text{cost}_0(\theta^T f^{(i)}) + \frac{1}{2} \sum_{j=1}^n \theta_j^2$$

- Now, we minimize using f as the feature vector instead of x
- By solving this minimization problem you get the parameters for your SVM
- In this setup, $m = n$
 - Because number of features is the number of training data examples we have
- One final mathematic detail (not crucial to understand)
 - If we ignore θ_0 then the following is true

$$\sum_{j=1}^n \theta_j^2 = \theta^T \theta$$

- What many implementations do is
 - $\theta^T M \theta$
 - Where the matrix M depends on the kernel you use
 - Gives a slightly different minimization - means we determine a rescaled version of θ
 - Allows more efficient computation, and scale to much bigger training sets
 - If you have a training set with 10 000 values, means you get 10 000 features
 - Solving for all these parameters can become expensive
 - So by adding this in we avoid a for loop and use a matrix multiplication algorithm instead
- You can apply kernels to other algorithms
 - But they tend to be very computationally expensive
 - But the SVM is far more efficient - so more practical
- Lots of good off the shelf software to minimize this function
- **SVM parameters (C)**
 - Bias and variance trade off

- Must choose C
 - C plays a role similar to $1/\text{LAMBDA}$ (where LAMBDA is the regularization parameter)
- Large C gives a hypothesis of **low bias high variance** --> overfitting
- Small C gives a hypothesis of **high bias low variance** --> underfitting
- **SVM parameters (σ^2)**
 - Parameter for calculating f values
 - Large σ^2 - f features vary more smoothly - higher bias, lower variance
 - Small σ^2 - f features vary abruptly - low bias, high variance

SVM - implementation and use

- So far spoken about SVM in a very abstract manner
- What do you need to do this
 - Use SVM software packages (e.g. liblinear, libsvm) to solve parameters θ
 - Need to specify
 - Choice of parameter C
 - Choice of kernel

Choosing a kernel

- We've looked at the **Gaussian kernel**
 - Need to define σ (σ^2)
 - Discussed σ^2
 - When would you choose a Gaussian?
 - If n is small and/or m is large
 - e.g. 2D training set that's large
 - If you're using a Gaussian kernel then you may need to implement the kernel function
 - e.g. a function
 - $f_i = \text{kernel}(x_1, x_2)$
 - Returns a real number
 - Some SVM packages will expect you to define kernel
 - Although, some SVM implementations include the Gaussian and a few others
 - Gaussian is probably most popular kernel
 - NB - make sure you perform **feature scaling** before using a Gaussian kernel
 - If you don't features with a large value will dominate the f value
- Could use no kernel - **linear kernel**
 - Predict $y = 1$ if $(\theta^T x) \geq 0$
 - So no f vector
 - Get a standard linear classifier
 - Why do this?
 - If n is large and m is small then
 - Lots of features, few examples
 - Not enough data - risk overfitting in a high dimensional feature-space
- Other choice of kernel
 - Linear and Gaussian are most common
 - Not all similarity functions you develop are valid kernels
 - Must satisfy **Mercer's Theorem**
 - SVM use numerical optimization tricks
 - Mean certain optimizations can be made, but they must follow the theorem
 - **Polynomial Kernel**
 - We measure the similarity of x and l by doing one of
 - $(x^T l)^2$
 - $(x^T l)^3$
 - $(x^T l + 1)^3$

- General form is
 - $(x^T l + \text{Con})^D$
- If they're similar then the inner product tends to be large
- Not used that often
- Two parameters
 - Degree of polynomial (D)
 - Number you add to l (Con)
- Usually performs worse than the Gaussian kernel
- Used when x and l are both non-negative
- **String kernel**
 - Used if input is text strings
 - Use for text classification
- **Chi-squared kernel**
- **Histogram intersection kernel**

Multi-class classification for SVM

- Many packages have built in multi-class classification packages
- Otherwise use one-vs all method
- Not a big issue

Logistic regression vs. SVM

- When should you use SVM and when is logistic regression more applicable
- If n (features) is large vs. m (training set)
 - e.g. text classification problem
 - Feature vector dimension is 10 000
 - Training set is 10 - 1000
 - Then use logistic regression or SVM with a linear kernel
- If n is small and m is intermediate
 - n = 1 - 1000
 - m = 10 - 10 000
 - Gaussian kernel is good
- If n is small and m is large
 - n = 1 - 1000
 - m = 50 000+
 - SVM will be slow to run with Gaussian kernel
 - In that case
 - Manually create or add more features
 - Use logistic regression or SVM with a linear kernel
- Logistic regression and SVM with a linear kernel are pretty similar
 - Do similar things
 - Get similar performance
- A lot of SVM's power is using different kernels to learn complex non-linear functions
- For all these regimes a well designed NN should work
 - But, for some of these problems a NN might be slower - SVM well implemented would be faster
- SVM has a convex optimization problem - so you get a global minimum
- It's not always clear how to choose an algorithm
 - Often more important to get enough data
 - Designing new features
 - Debugging the algorithm
- SVM is widely perceived a very powerful learning algorithm